

Using Dynamic Libraries and Plugins in C++

Thomas Witt
witt@styleadvisor.com

April 20. 2006

Welcome to the
basement of software
construction

Dynamic Shared Object

- Various names
 - Dynamic Library
 - DLL
 - Shared Library
 - Plugin
 - Module
- A means of packaging executable code
 - More executable than library

Dynamic Linking

- Defer final program assembly until load time
- Benefits:
 - Shared Libraries
 - Delay loading
 - Explicit loading
 - Easily serviceable libraries (well...)
- Costs:
 - Increased startup time

What the standard says

...

Terminology

- Translation unit (C++03)
 - Preprocessed source file
- Compiled Translation Unit
 - Object file
- Symbol
 - Named entity in an compiled translation unit
- Load unit
 - Executable, dynamic library

Terminology (Cont.)

- Direct dependencies of a load unit
 - Other load units available at static linktime
- Dependencies of a load unit
 - Transitive closure of direct dependencies
- Load set
 - Primary load unit and its dependencies

Compilation Model

- The compiler operates on individual translation units

- The static linker operates on set of compiled translation units and their direct dependencies
- The dynamic linker operates on load sets

Linker/Loader

- Operating system facility
- Task: Create a runnable program in memory from compiled translation units
- Linker
 - Symbol binding
 - Relocation
- Loader
 - Program loading
 - Relocation

Object File Format

- Linker meta information
 - Header
 - Relocations
 - Symbols
 - Runtime dependencies
- Code
- Text (data)
- Debugging information

Object File Format

- Common Object File Format (COFF)
System V, AIX (XCOFF)
- Portable Executable Format (PE)
Windows, BeOS R3
- Executable and Linking Format (ELF)
System V R4, Solaris, IRIX, BSD, Linux, BeOS
- Mach object file format (Mach-O)
Mac OS X, Darwin, NeXTSTEP

Symbol Visibility

- Visibility at load unit level
 - Distinct from linkage that works per translation unit (static anybody?)
- External visibility
 - Symbols are visible and accessible from other load units
- Internal visibility
 - Symbols are neither visible nor accessible from other load units

Symbol Visibility (PE)

- Internal visibility by default
- Export symbol tables
- `__declspec(dllexport)`
 - Marks symbols as externally visible
- `__declspec(dllimport)`
 - Improves code generation for imported symbols

Symbol Visibility (ELF)

- External visibility by default
- `__attribute__((visibility(option)))`
 - default → external visibility
 - hidden → internal visibility
 - internal
 - protected
- `#pragma`
- Compiler switch
- Symbol tables

Symbol Lookup (PE)

- Symbols are bound to load units at static link time
 - Import libraries
 - Load units can forward symbols
- Symbols are resolved on first use
- Lookup by
 - Hint and name
 - Ordinal



Symbol Lookup (ELF)

- Symbols are bound to load units at dynamic link time by default
 - Symbol interposition
- Symbols can be bound to load units at static link time
 - Direct binding
 - Interposers
- Lookup by name
 - Hash table


Symbol Lookup (ELF)

- Search scope
 - world
 - group
- Visibility (ELF)
 - world
 - local
- Load units are searched in load order

Symbol Lookup (ELF)



Symbol Lookup (ELF)



Symbol Lookup (Mach-O)

- Symbols are bound to load units at static link time by default
 - Two level namespace
 - Flat Namespace
- Otherwise like ELF

C++

finally

Issues

- DSO compatibility
 - ABI configuration
 - ABI stability
 - Versioning
- Packaging
 - Encapsulation
 - Symbol resolution

ABI

- Application Binary Interface
 - Interface for compiled object code
- Components
 - Memory layout
 - Function call interface
 - Compiler generated objects
 - Name mangling
- Itanium C++ ABI 58 pages
 - C++ Standard Library not covered

Name Mangling

- Mapping C++ names to symbol names
 - Overloaded functions
 - Templates
 - Namespaces
- `namespace bar { int foo(double); }`
 - Microsoft
 - `?foo@bar@@YAXN@Z`
 - GCC 4.0.1
 - `_ZN3bar3fooEd`

Extern "C"

- extern "C" gives C-Linkage to a declaration
 - Simpler C name mangling rules
 - C calling conventions
- C-Linkage reduces the available feature set.

ABI Stability

- Rules for dynamic library APIs according to Apple (C++ Runtime Environment Programming Guide 2005-08-11)
 - Avoid inheriting from Classes in the standard C++ library
 - In your class definition, avoid member variables whose type is a class defined by the C++ runtime
 - Avoid using standard C++ runtime classes for arguments and return values.

Compiler Generated

- Class members
- Inlines
- Template Instantiation
 - Extern Template
- Run Time Type Information (RTTI)
- Virtual Function Table
- Compiler generated objects do not have a clear binding to a single translation unit
- Collapsing requires load set analysis

RTTI

- Application
 - typeid()
 - Exception handling
 - dynamic_cast (polymorphic types)
- RTTI might be emitted in multiple object files

More Possible Duplicates

- Global operator new
 - Requires symbol interposition
- Static objects
 - Templates
 - Inlines
- Inlines
 - Address of


Modules

- Load semantics
 - Thread local storage (PE)
- ABI stability
 - Modules should work across application releases
- Encapsulation
 - Name clashes
 - Memory management

Programmatic Loading

- Loading load units at runtime
 - dlopen, dlsym
 - ELF
 - Mach-O
 - LoadLibrary, GetProcAddress
 - PE
- Modules
- Plug-Ins

Symbol Lookup (ELF)



Modules

- C-Interface
- Control memory allocation
 - Do not expose allocated memory
 - Provide deallocation API
- Link statically
 - Alas, some libraries are dynamic only

Using C++ to implement modules today is unsafe in some scenarios

Application Startup

- Relocations

- Preferred load address (PE, ELF)
- Position independent code (ELF, Mach-O)

- Symbol binding

- Reduce symbol table size
- Shorten symbol names (ELF, Mach-O ?)

Summary

- Dynamic shared objects form integral part of modern program design
- C++03 is oblivious to DSOs
- C++ is not very well suited for packaging in independently compiled DSOs
- Using DSOs may require intimate platform knowledge
- Poor portability
- C++0x may (shall ?) address DSOs

References

- John R. Levine
"Linkers and Loaders"
Morgan-Kaufman, Oct 1999, ISBN 1-55860-496-0
- Ulrich Drepper
"How To Write Shared Libraries"
Jan 22 2005
- "Itanium C++ ABI (\$Revision: 1.86\$)"
<http://www.codesourcery.com/cxx-abi/abi.html>
- "Linker and Libraries Guide"
Sun Microsystems, Inc., December 2003

References

- Benjamin Kosnik
“Dynamic Shared Objects: Survey and Issues”
Nov 2 2005
- Matt Austern
“N1400 Toward standardization of dynamic libraries”
Sep 25 2002
- “Mac OS X ABI Mach-O File Format Reference”
Apple Computer, Inc., Mar 8 2006
- “Mach-O Programming Topics”
Apple Computer, Inc., Nov 9 2005
- “Dynamic Library Programming Topics”
Apple Computer, Inc., Feb 7 2006

References

- Russ Osterlund
"What Goes On Inside Windows 2000: Solving the Mysteries of the Loader"
- Mark Pietrek
"An In-Depth Look into the Win32 Portable Executable File Format"
- Mark Pietrek
"An In-Depth Look into the Win32 Portable Executable File Format, Part 2"
- Mark Pietrek
"Optimizing DLL Load Time Performance"
- Kang Seonghoon
"Microsoft C++ Name Mangling Scheme"
Version 1.1, Nov 25 2005